

Chapter 3: Reading the Room & Red Flags

The Morning After the Call

It's 7:00 AM. Your coffee is cold.

In front of you is the legal pad from yesterday's marathon of discovery calls. You stare at the page, and the page—a chaotic mess of frantic scribbles—stares back.

- Client A: "Gmail/Notion/Slack. Broke. Budget: \$2k MAX (incl. me??)"
- Client B: "15 PAID AI TOOLS?! Using 20%... Lacking all direction."
- Client C: "Salesforce + HubSpot + Zoho (why?!). Needs 'shiny things' for board. 'Spaghetti back-end'."
- And scrawled in the margin, circled in angry red ink: "Needs AI agents NOW. Saw on LinkedIn."

This... *this* is the job.

Yesterday was "Day 2: The Art of the Discovery Call".¹ That was data collection. This morning, with your cold coffee, is "Day 3." This is the triage. This is *part two*.¹

Discovery calls fail because of noise, and it's coming from both sides of the phone.¹ You were rambling, "adding way too much context".¹ They were confused, "pivoting left/right/center".¹ The message was lost in translation.

This chapter is the bridge. It's the mental framework for turning that call's chaos into a clear diagnosis. It's how you cut through the noise when articulation fails on one side and digestion fails on the other.¹

Let's learn how to read the room.

Visual Suggestion: A high-contrast, top-down photo of the desk described above. A white legal pad is covered in chaotic, handwritten notes, perfectly capturing the "noise" of a

discovery call. A half-empty, cold coffee mug sits beside it.

The Triage: Pattern-Matching Your Client

After enough calls, you'll develop a consultant's sixth sense. It's simple "pattern matching".¹ You'll be able to diagnose a client in the first 90 seconds, just like a seasoned doctor.

You're going to meet three "patients" every single day. Let's get to know them.

The Bootstrapper (Scenario A: 5-20 People)

This is the 15-person team, the "real startup, real broke, maybe bootstrapped" crew.¹ Everyone uses the free version of ChatGPT. Their entire world runs on the simple stuff: Gmail, Notion, and Slack.¹

The Diagnosis (The Math):

They have a budget of "\$500-\$2K/month".¹ Here's the catch: that includes your fee.¹ They are "very risk-averse".¹ If their max is \$2K, their real AI budget is probably \$1K.

So, what's your prescription?

It's counter-intuitive. They're broke, so you're going to tell them to spend money.

The Prescription:

"Just pay for ChatGPT Teams".¹

Let's run the numbers so you can explain it. At \$25/person/month, a 20-person team costs \$500/month.¹ This will be 25% to 50% of their *entire* AI budget.¹ You must justify this.

Here's how: "Can you do everything with just ChatGPT? Yes".¹ Why? Because their simple stack—Gmail, Notion, Slack—all have "integrations, connectors, MCP servers".¹ Everything can route through ChatGPT Teams. This one move gives them:

1. Higher security than the free version (a huge win).¹
2. The ability to use agent builders, custom GPTs, projects, and workflows.¹
3. A path to connect to any API with custom schemas, even without a native connector.¹

This team needs **ONE tool done right, not 15 tools done poorly**.¹ Your value here isn't in

building some complex, custom mess. Your value is in *consolidation*.

Visual Suggestion: A simple, stark pie chart titled "The Bootstrapper's AI Budget (\$1,000/mo)". A large 50% slice is labeled "ChatGPT Teams (\$500)" to visually hammer home the budget impact and the 'one tool' strategy.

The Overwhelmed Director (Scenario B: 21-100 People)

This patient is different. They are "Not broke. Not lacking resources. Lacking direction".¹

This client has 10-15 paid AI tools but is "using each tool only 20% of potential".¹ They didn't even know ChatGPT could make images like Midjourney. Their knowledge is in "multiple places"¹, and their budget is a healthy \$2K-\$10K/month.¹

The "Foot in the Door" Play:

Your goal here is not a big project. Your goal is a "small retainer to keep your foot in the door".¹ Think \$200-\$500/month.

This is a classic "land and expand" strategy. The \$500/month isn't the product; trust is. You're not selling a service; you're selling a low-cost, multi-month trial for a future \$50,000 strategy contract. You stay "embedded".¹ You become the "trusted advisor".¹ When the *real* workshop or roadmap opportunity comes up, you're already there.

The Phased Rollout:

You'll use your retainer to guide a phased rollout 1:

1. **Start Small:** Grab 10 "power users" in "revenue-adjacent roles" (like the sales team).¹
2. **Run a Pilot:** Run a 2-3 week pilot to get quick wins.¹
3. **Centralize:** Get their "scattered data" into one central repo (Google Drive, Dropbox, whatever).¹
4. **Consolidate:** Cut those 15 tools down to 5, or even 3.¹

The "Secret Weapon":

During this pilot, you'll introduce them to your secret weapon: Claude Code.¹ This is your "Swiss Army knife".¹ You'll sit them down and say:

"I don't care if you're not developers. I don't care if you're afraid of the terminal. This is what you can do. This is your magic Swiss Army knife for building anything at scale".¹

Focus on the sales department first. Why? Because if they report to a board, they need a "tangible return on initial value".¹ When they get that "first taste of what's possible," more

budget gets unlocked, and the skeptics start paying attention.¹

The "Blue Sky" Session (The Real Magic)

After your 2-3 week pilot is a success, you run this session. This is the payoff.

You bring in the 10 power users who just experienced their "first aha moment".¹ You set the stage, and then you ask them three questions:

1. *"Now that you know just a sliver of what's possible, what do you think would change your day-to-day life?"¹*
2. *"How could you become 10x more efficient?"¹*
3. *"In a perfect world where AI didn't hallucinate, was affordable... what would you want to accomplish with it?"¹*

Will you grant 90% of these wishes? No.¹ But you will hear incredible ideas from people who would never otherwise speak up.¹

Now you have a pile of dreams. Your job is to turn them into a systematic roadmap. You do this with the **RICE scoring framework**.¹

Visual Suggestion: A clean infographic illustrating the RICE formula. A simple equation on a whiteboard: $(\text{Reach} \times \text{Impact} \times \text{Confidence}) / \text{Effort} = \text{Priority Score}$.

You take every idea and score it¹:

- **Reach:** How many people are affected?
- **Impact:** What gets solved? (How big is the impact?)
- **Confidence:** Out of 100, how confident are you that you can *actually* build this?
- **Effort:** How much work is required?

This systematic process allows you to prioritize their wishes. You can now go back to the C-suite with a data-driven plan, like this:

Table 1: Prioritizing the Blue Sky Ideas

Blue Sky Idea	Reach (How many)	Impact (What gets)	Confidence (Can we)	Effort (How)	RICE Score (ER×I×C)

	people?)	solved?)	build this?)	much work?)	
1. Automate Sales Reports	10 (Sales Team)	8 (Saves 20 hrs/wk)	9 (High)	3 (Low)	240 (High Priority)
2. AI Social Media Agent	3 (Marketing)	5 (Med Impact)	4 (Low)	8 (High)	7.5 (Low Priority)
3. Internal Knowledge Bot	50 (Whole Co.)	6 (Med Impact)	7 (Med-High)	7 (Med-High)	30 (Mid Priority)

The Enterprise Labyrinth (Scenario C: 100+ People)

This is the final patient. The Enterprise. This client has "8 systems where 1 would satisfy a normal company".¹ We're talking Workday, Greenhouse, Salesforce, and HubSpot.¹ Their knowledge is "heavily siloed." Their approvals are pure "bureaucracy, red tape".¹

The Brutal Reality:

You are not going to fix their system architecture. That's a "6-12 month" job.¹ These enterprises, especially if they have public or private shareholders, need to LOOK like they're embracing AI. They need "shiny things that work NOW".¹

The Hack (The Quick Win):

Your hack is to build something that lives OUTSIDE their spaghetti linguine tortellini mess of the back end.¹

Here's an example: Their sales team struggles to remember their "522 product SKUs".¹ You "vibe code" a simple knowledge-sharing tool. It's one central repo, immediately accessible, populated by user research.¹

The tool itself isn't the point. The point is to "**show the org what 'moving quick' looks like**".¹ This single, fast win can "break down red tape ideologies" and "inspire them to throw resources at manifesting the AI dream".¹

Enterprise Warning: With these clients, it is "**never the same content twice**".¹ Every

situation is unique. You'll face "political factors, EU regulations, medical constraints, controversial products," and more.¹ Be prepared.

The Prescription: Being the "Politely Savage" Doctor

You've triaged the patient. You've made your diagnosis. Now, it's time to write the prescription.

This is where you must be an "AI Doctor, not an AI drug dealer".¹ The drug dealer gives the client what they *want* (the sexiest new tool). The doctor gives them what they *need* (the foundations).

The Doctor's Bag (Choosing the Platform)

Your recommendation framework is your medical bag. You choose the right tool for the *actual* ailment¹:

- **Prescribe ChatGPT when...** they need "broad integrations," are "Google Workspace heavy," need "team collaboration," and want "one ecosystem" (plus image/voice features).¹
- **Prescribe Claude when...** they need "deep reasoning," "complex analysis," and have "document-heavy workflows" (as long as the docs aren't too large).¹
- **Prescribe Gemini when...** they are "Google ecosystem locked in," need "video input" (like YouTube integration), or are extremely "cost-sensitive" (thanks to its "gracious free rate limits").¹
- **Prescribe a Custom Build when...** it is an "enterprise play only," they have the resources, and the "Blue sky scenario made it obvious".¹

The "Politely Savage" Truth (Explaining Trade-offs)

This is the hard part of the consultation. Your recommendation *must* "come with sacrifice".¹

They'll have "touchy-feely" objections.¹ "But we love the yellow and blue design in Gamma!".¹

Your recommendation, Claude, makes PowerPoints that look "more dry".¹

You have to rip these off like a bandaid.¹ It's time to be "politely savage".¹

Here is your script. Memorize it.

"Like a personal trainer, being nice when you're 100 pounds over your weight loss goal won't help. It's not gonna help you. You need to cut refined carbs and hit the weights 5x a week."¹

You are not their friend; you are their partner in getting to a goal. This builds camaraderie. You're being strict because you want them to *win*. This builds trust. And that trust lets you make the "honest advisor play."

This is, perhaps, the most important monologue in this entire playbook.

"I could charge you hourly and you'd need me for 6 months. I could accumulate billables. But I want to sleep at night knowing I gave you highest information density per minute. I want to get you there in as few steps as possible."¹

That single speech is what separates you from every "drug dealer" they've ever met.

Visual Suggestion: A two-panel cartoon. Panel 1: An "AI Drug Dealer" in a trench coat, offering 15 shiny, colorful bottles ("Midjourney," "Gamma," "Agents!"). Panel 2: An "AI Doctor" in a lab coat, holding a single, boring, professional bottle labeled "Consolidated Stack."

The Redundancy Plan (Always Have a Plan B)

Your 30/60/90-day roadmap is Plan A. But every organization has its own "heartbeat".¹ Plan A might fail.

You need a Plan B for when you get "crazy C-suite pushback" or "Luddites poison-pill the company".¹ What's the backup?

- Open source models?
- Self-hosted N8N instances?
- A "no licenses needed" plan if someone blocks all subscriptions? ¹

You need this emotional armor, because here's the "real talk"¹:

"I've been on calls with 10 people who hate my existence just because I'm trying to help

*them embrace AI. You won't always be seen as the hero. Sometimes you're the villain."*¹

The Exit Interview: A Final Checklist of Red Flags

You've delivered the diagnosis and the prescription. This is your final act of value: the discharge papers. You leave them with a warning list to protect them after you're gone.

The Patient's Discharge Papers (Common Pitfalls)

Here is the list of pitfalls you must warn them about ¹:

- **Don't chase features.**
- **Don't overbuild** (e.g., "don't use 10 agents because YouTube made it look sexy").¹
- **Don't skip foundations.**
- **Don't rush rollouts.**
- **Don't forget security** (some companies "will streak naked with customer data").¹
- **Don't ignore training.**

That last one is special. It isn't just a warning; it's your single greatest upsell opportunity.¹ When you tell them "Don't ignore training," you are creating the need for the very workshop you plan to sell them next. It's the most ethical sales pitch there is.

When to Pump the Brakes (The Red Flags)

Finally, you tell them about the clients you avoid. This shows you have standards and reinforces your "honest advisor" status. These are the red flags that scream "this client will be a nightmare".¹

"We need AI agents NOW."

- **Your Diagnosis:** "No you don't. You need to get your grip together. Even in the best workflows, AI agents are used once or twice max. Determinism is the best way to build a workflow".¹

"I saw this on LinkedIn."

- **Your Diagnosis:** "Scope creep superhero. Red flag incarnate".¹

"My competitor has it."

- **Your Diagnosis:** "You're building AI for wrong reasons. Different companies, different heartbeats, different bureaucracy. What works for them won't work for you".¹

"It MUST be cutting-edge."

- **Your Diagnosis:** "Cutting edge = bleeding edge = high risk of bugs, landscape changes, technology shifts. Don't build enterprise-scale on something that came out last week".¹

When you see multiple red flags, you "pump the brakes. These clients will be more headache than blessing".¹

The Future-Proof Promise

You close the session by bringing it all home. Your goal isn't to sell them bells and whistles. Your goal is to give them the "minimum stack this organization needs".¹ You give them a plan they can stick to, "no matter what Gemini or Anthropic releases next month".¹

That is how you transcend the noise. That is how you become the trusted advisor.

In the next one, we'll look at solution design—how to architect the AI implementations that actually stick.¹

Works cited

1. Consulting Playbook Framework.xlsx